



Home of Redis

Redis Enterprise Developer Workshop

Modules Overview



Redis Labs – Home of Redis



The commercial company behind Open Source Redis




Provider of the **Redis Enterprise (Redis^e)** technology, platform and products

Founded in 2011

HQ in Mountain View CA, R&D center in Tel-Aviv IL

Topics



- Overview of the Redis Module System
- Redis-ML and Spark-ML modules
- reJSON
- Redisearch
- Redis Graph



Overview of Redis Modules

Redis Modules


- Any C/C++ program can now run on Redis
- Use existing or add new data-structures
- Enjoy simplicity, infinite scalability and high availability while keeping the native speed of Redis
- Can be created by anyone



Modules : A Revolutionary Approach

Adapt your database to your data, not the other way around

Neural Redis Simple Neural Network Native to Redis	Redis-ML Machine Learning Model Serving	RediSearch Full Text Search Engine in Redis
ReJSON JSON Engine on Redis. Pre-released	Time Series Time series values aggregation in Redis	Graph Graph database on Redis based on Cypher language
Rate Limiter Based on Generic Cell Rate Algorithm (GCRA)	Crypto Engine Wrapper Secure way to store data in Redis via encrypt/decrypt with various Themis primitives	Secondary Index/RQL Indexing + SQL-like syntax for querying indexes. Pre-released




Redis-ML and Spark-ML



Redis-ML Module

- Predictive models as a native Redis Data Type (4.0 Modules API)
- Predictive model serving engine
- Multiple supported models
- Training framework independent

Redis ML Module



Redis Module

Tree Ensembles

Linear Regression

Logistic Regression

Matrix + Vector Operations


More to come...

Redis-Spark Module

- Connector between Apache Spark and Redis-ML
- Train models using Apache Spark
- Load trained models into Redis

Example: Random Forest

- A collection of decision trees
- Splitter Node can be:
 - Categorical (e.g. day == "Sunday")
 - Numerical (e.g. age < 43)
- Decision is taken by the majority of decision trees



Forest Data Type Commands

Add nodes to a tree in a forest:

```
ML.FOREST.ADD <forestId> <treeId> <path>  
  [ [NUMERIC|CATEGORIC] <splitterAttr> <splitterVal> ] |  
  [LEAF] <predVal>
```

Perform classification/regression of a feature vector:

```
ML.FOREST.RUN <forestId> <features>  
  [CLASSIFICATION|REGRESSION]
```

Forest Data Type Example

```
> ML.FOREST.ADD myforest 0 . CATEGORIC sex "male" .L  
LEAF 1 .R LEAF 0
```

OK

```
> ML.FOREST.RUN myforest sex:male  
"1"
```

```
> ML.FOREST.RUN myforest sex:something_else  
"0"
```

Real World Challenge: Ad Serving

- Need to serve 20,000 ads/sec @ 50msec data-center latency
- Runs 1k campaigns → 1K random forest
- Each forest has 15K trees
- On average each tree has 7 levels (depth)

Real World Challenge

- Ad serving company
- Need to serve 20,000 ads/sec @ 50msec data-center latency
- Runs 1k campaigns → 1K random forest
- Each forest has 15K trees
- On average each tree has 7 levels (depth)

Ad Model Serving: Infrastructure Cost Saving

Cut computing infrastructure

by 97%




Homegrown

1,247 x c4.8xlarge



35 x c4.8xlarge



reJSON



What is JSON

- JavaScript Object Notion
- Data interchange format
 - Lightweight
 - Human Readable
 - Simple to Parse
 - Simple to Generate
- ECMA Standard
- Common web development format

```
{  
  "userId": 0,  
  "firstName": "Tague",  
  "lastName": "Griffith",  
  "userName": "tague"  
}
```

Storing JSON in Redis

Serialized as Redis String

- GET/SET operations
- $O(N)$ access
- Client deserialization
- No in-place updates

Deserialized in Hash

- Decode: HMSET
- Encode: HMGET
- $O(1)$ Access
- Client deserialization
- No in-place updates

reJSON Module

- JSON as a native Redis Data Type (4.0 Modules API)
- Keys map to JSON values
 - Scalars
 - Objects
 - Arrays
 - Nested or Not
- Stored as a document tree structure
- Path access to JSON elements
- Atomic, in-place updates

reJSON Basic Get/Set Operations

```
redis> JSON.SET json:scalar . '"Hello JSON!'"
```

```
OK
```

```
redis> JSON.SET json:object . '{"userId": 3001, "firstName": "Tague"}'
```

```
OK
```

```
redis> JSON.GET json:scalar
```

```
"Hello JSON!"
```

```
redis> JSON.GET json:object
```

```
{"userId":3001,"firstName":"Tague"}
```

```
redis> JSON.GET json:object .userId
```

```
3001
```

```
redis> JSON.SET json:object .userId 2001
```

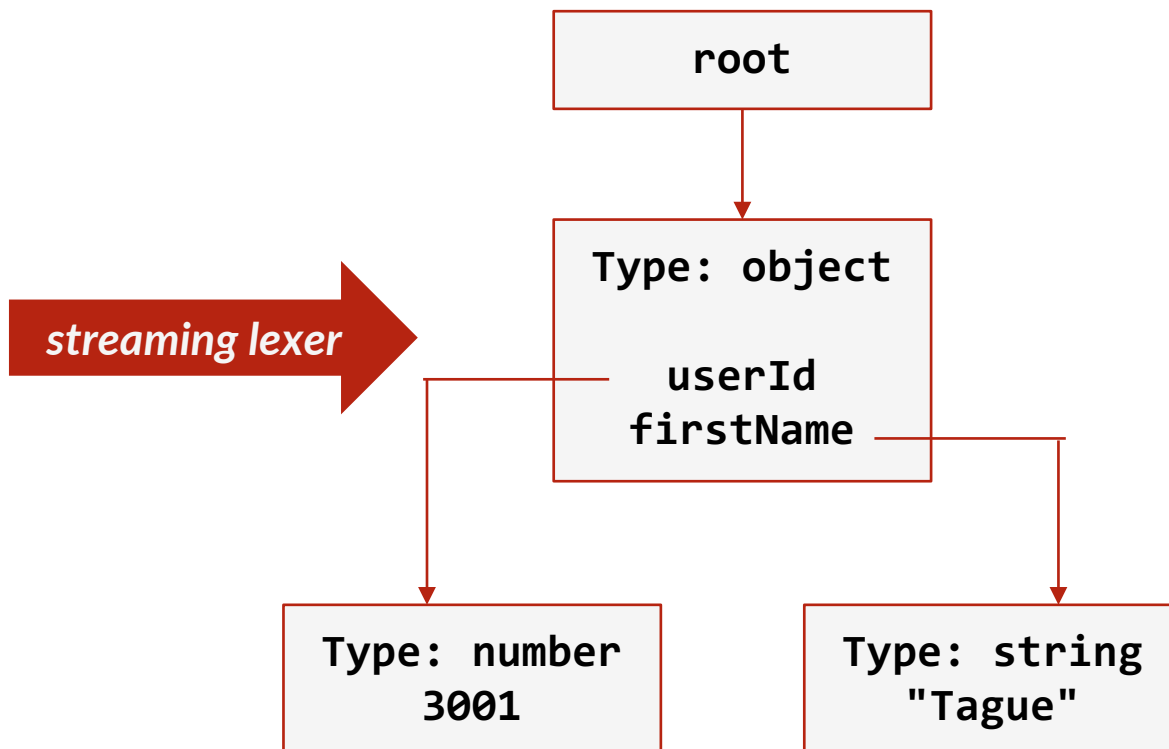
```
OK
```

```
Redis> JSON.GET json:object
```

```
{"userId":2001,"firstName":"Tague"}
```

ReJSON tree data structure

```
JSON.SET key .  
{  
  "userId": 3001,  
  "firstName": "Tague"  
}
```




reJSON Commands

General	JSON.DEL, JSON.GET, JSON.MGET, JSON.SET, JSON.TYPE
Numbers	JSON.NUMINCRBY, JSON.NUMMULTBY
Strings	JSON.STRAPPEND, JSON.STRLEN
Objects	JSON.OBJKEYS, JSON.OBJLEN
Arrays	JSON.ARRAPPEND, JSON.ARRINDEX, JSON.ARRINSERT, JSON.ARRLEN, JSON.ARRPOP, JSON.ARRTRIM, JSON.RESP


Access reJSON API


- Python
 - rejson-py
 - redis-py execute command
- Java
 - JReJSON
 - JedisConnection.execute

Throughput



Latency





Redisearch



Redisearch Module

- Developed by Redis Labs using 4.0 Module API
- Optimized Data structures
- Fast index and search
- Text, numeric and geo filters
- Non-blocking updates and deletes
- Scale to billions of documents

Redisearch Fundamentals

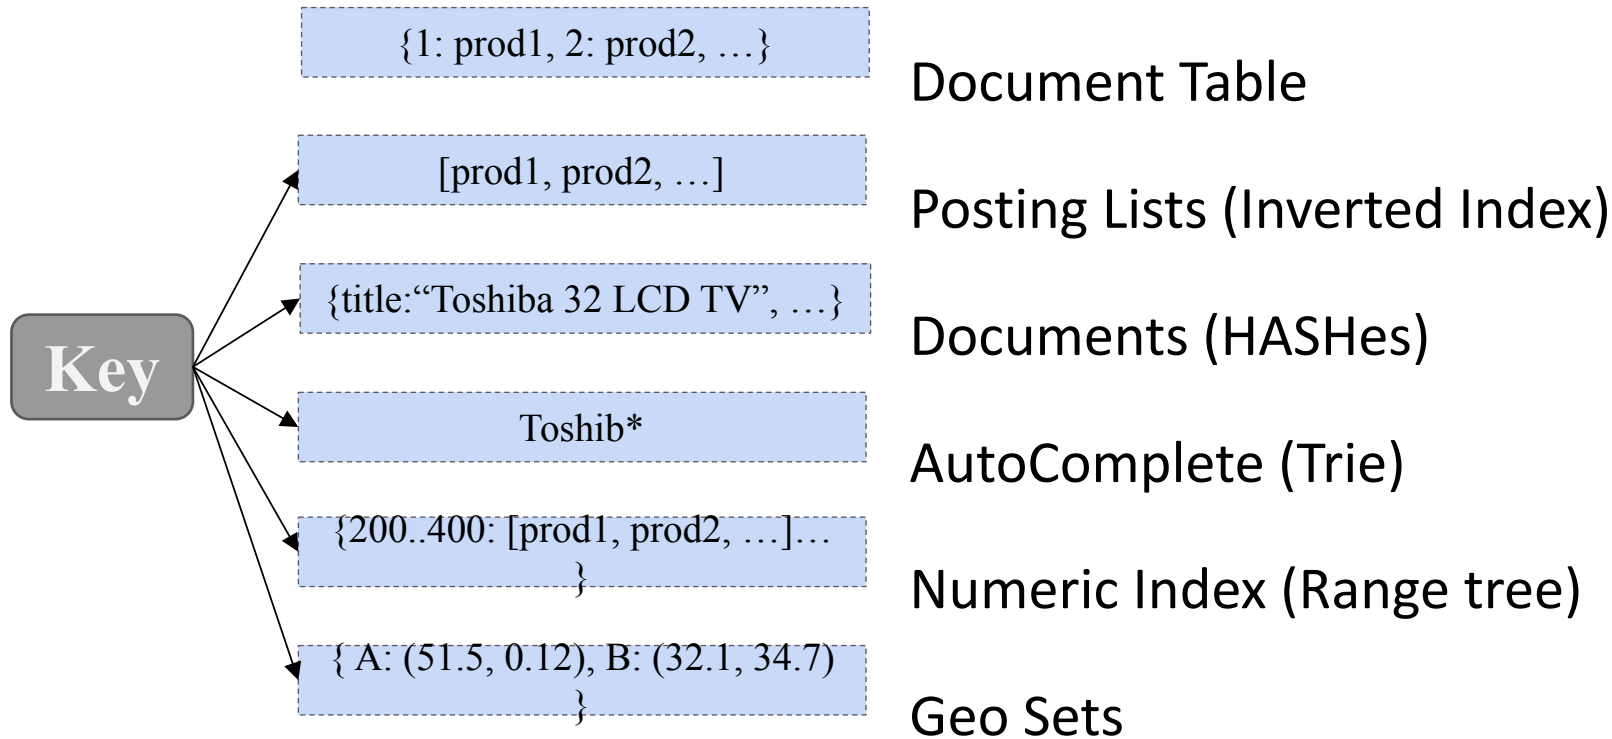
```
> FT.CREATE products SCHEMA title TEXT price NUMERIC

> FT.ADD products prod1 1.0 FIELDS title "Toshiba 32 LCD TV" price 250
> FT.ADD products prod2 1.0 FIELDS title "Samsung 42 LCD TV" price 350

> FT.SEARCH products "lcd tv"
1) (integer) 2
2) "prod1"
   ...
4) "prod2"
   ...

> FT.SEARCH products "lcd tv" FILTER price 300 400
1) (integer) 1
"prod2"
```

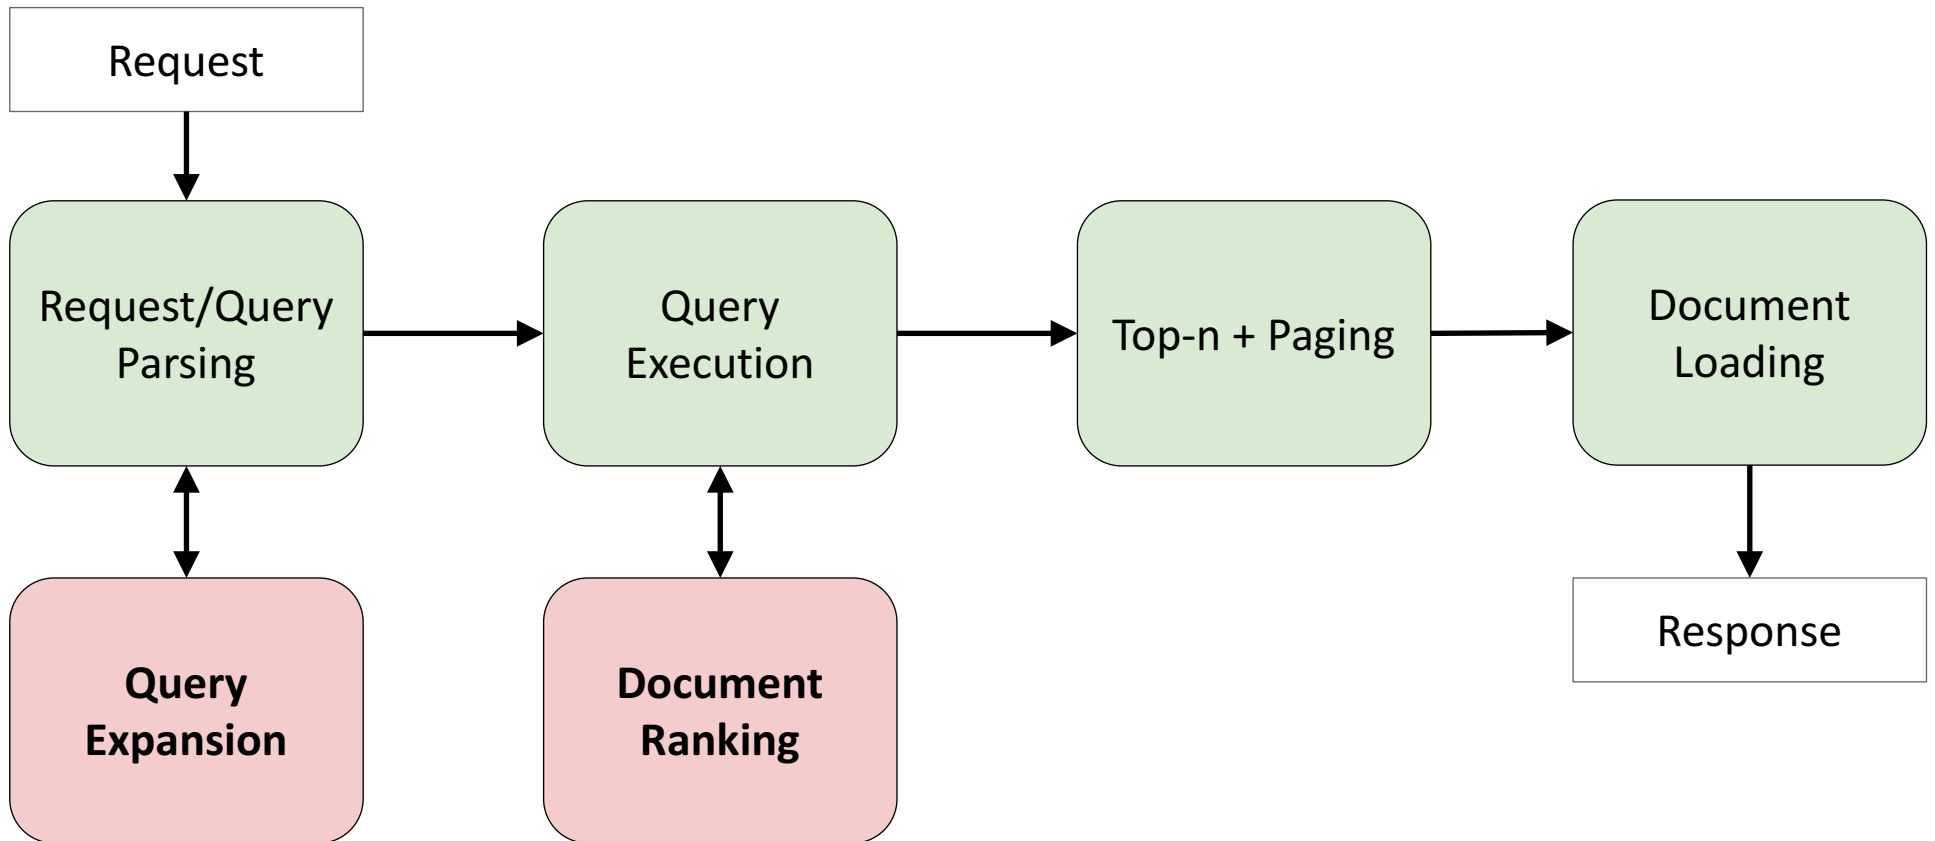
Redisearch Data Structures



Redisearch Query

```
FT.SEARCH {index} {query} [NOCONTENT] [VERBATIM]
[NOSTOPWORDS] [WITHSCORES] [WITHPAYLOADS]
[FILTER {numeric_field} {min} {max}] ...
[GEOFILTER {geo_field} {lon} {lat} {radius} m|km|mi|ft]
[INKEYS {num} {key} ... ]
[INFIELDS {num} {field} ... ]
[SLOP {slop}] [INORDER]
[LANGUAGE {language}]
[EXPANDER {expander}]
[SCORER {scorer}]
[PAYLOAD {payload}]
[SORTBY {field} [ASC|DESC]]
[LIMIT offset num]
```

Redisearch Search Request Processing



Evaluating Queries

- Document at-a-time approach
- Indexes are decoded as iterators, one object per round
- Top N results copied into a heap & sorted on the fly

Scoring Results

- Each document in the query's result is evaluated by a **Scoring Function**
- **Top-N** scored documents are selected
- The default uses:
 - Term Frequency–Inverse Document Frequency (TF-IDF)
 - User given a priori document score
 - Proximity boosting
- And you can write your own function!

Auto Complete

- Custom data type - trie based
- Fuzzy completions
- Manual dictionary building and scoring
- Unicode aware case-folding

Scaling RediSearch

- Partition by document id
- Each partition contains a full index of $1/N$ of the documents
- All terms for a sub-index are stored on the same shard
- This means we need a query coordinator:
 - Distribute the queries to all shards
 - Merge the results

Redisearch Sharing



Redis

+



Search Engine

+



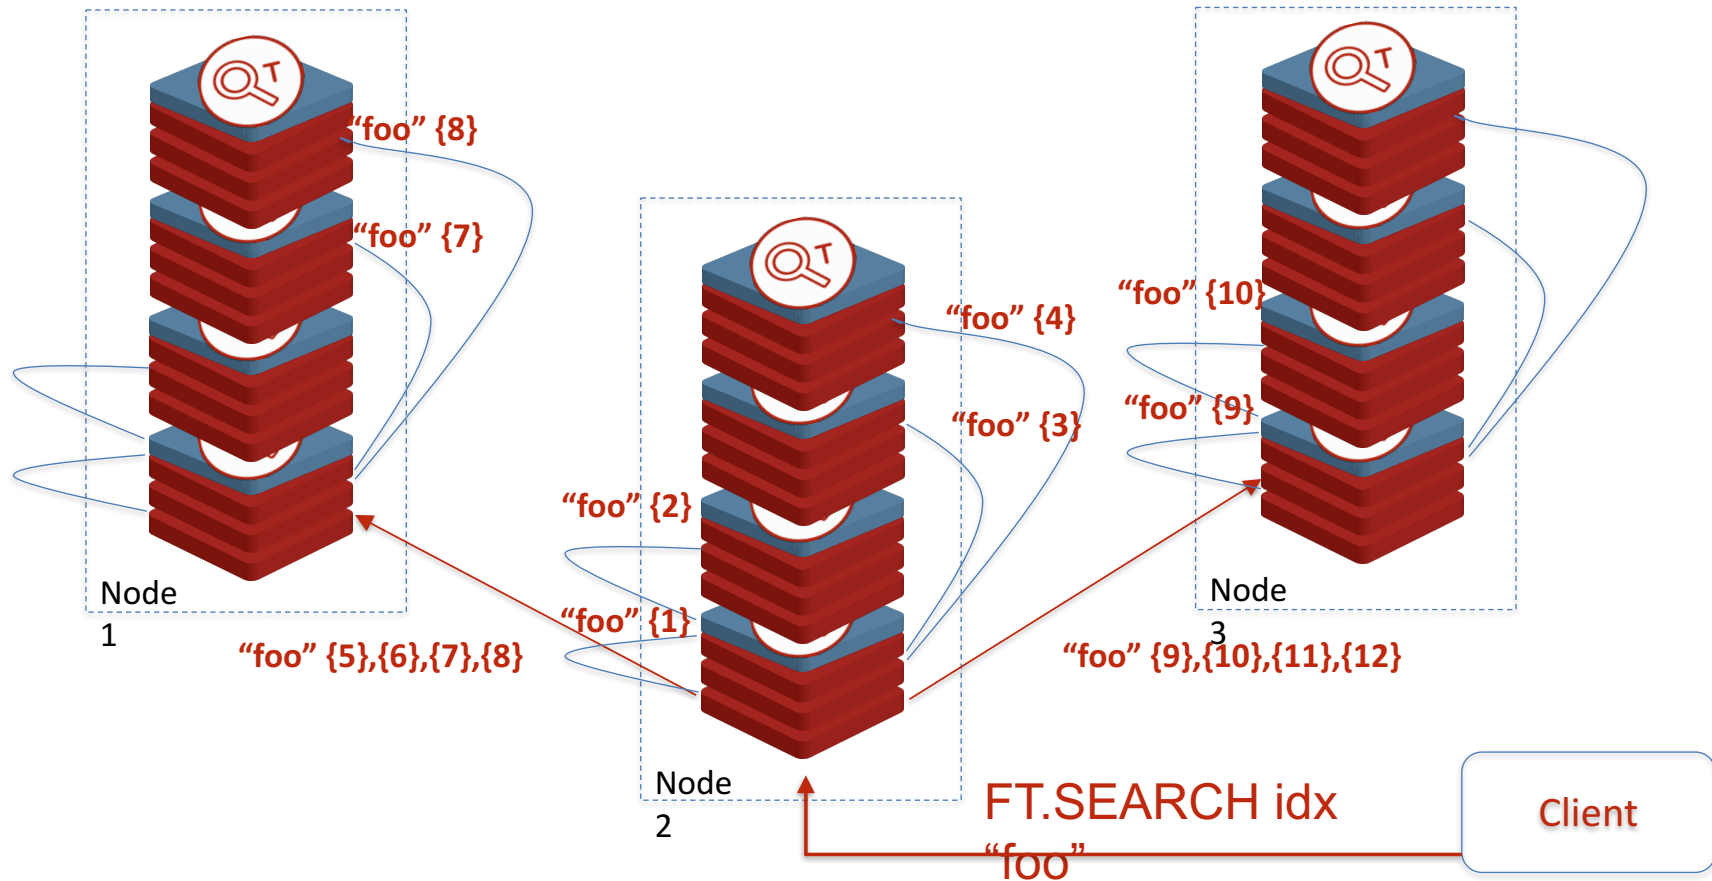
Coordinator

=





Redisearch Shard

RediSearch Cluster Architecture



Performance

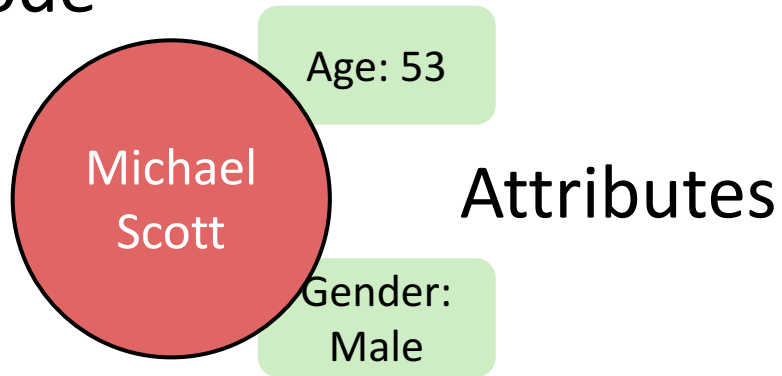





Redis Graph

Basic Graph Data Model

Node



Relation



Adding Data

```
GRAPH.CREATENODE <Graph_Name> <Label> <Attributes>
```

```
GRAPH.CREATENODE the_office employee name "michael scott" age 53 gender male  
status single
```

```
GRAPH.ADDEDGE <Graph_Name> <SRC_NODE> <Relation> <DEST_NODE> <Attributes>
```

```
GRAPH.ADDEDGE the_office michael scott boss jim halpert <Attributes>
```

HexaStore Storage Format

S


Subject

P

Predicate


O

Object



SPO OSP
SOP PSO
OPS POS

HexaStore Example



Triplets

SPO:Michael:Boss:Jim
SOP:Michael:Jim:Boss
OPS:Jim:Boss:Michael
OSP:Jim:Michael:Boss
PSO:Boss:Michael:Jim
POS:Boss:Jim:Michael

Converting Queries to Predicates

Who does Michael manages?


SPO:Michael:Boss:*

Who manages Jim?

OPS:Jim:Boss:*

Who manages who?

PSO:Boss:*




Query Language Cypher

```
MATCH <Graph_name> <graph_pattern>
WHERE <filters>
RETURN <entities>
ORDER BY <properties>
LIMIT <count>
```

Cypher Example

```
MATCH (Michael:employee {name:michael  
scott})-[boss]->(E:employee)-[visit]->  
(c:country)  
WHERE (E.age > Michael.age AND c.continent  
= Europe)  
RETURN E.name, count(c.name) AS  
countriesVisited  
ORDER BY countriesVisited, E.age DESC  
LIMIT 5
```

Query Execution



Benchmark

- 150K inserts per second
- 15K queries per second

redislabs

Home of Redis

