



EVERYWHERE.

Transactions and Concurrency Control in Redis

RedisConf18 Introductory Training



#Redis #RedisConf

presented by:
redislabs
home of redis

Agenda

- What are transactions
- Redis Execution Model
- Transaction Commands: MULTI, EXEC, DISCARD
- Optimistic Concurrent Control: WATCH, UNWATCH
- Durability
- Pipelines



ACID Transactions

redislabs

ACID Transactions

A - Atomicity

- Transaction executes as an indivisible unit

C - Consistency

- Transaction takes database from one valid state to another

I – Isolation

- Transactions result in a state as if they were executed sequentially

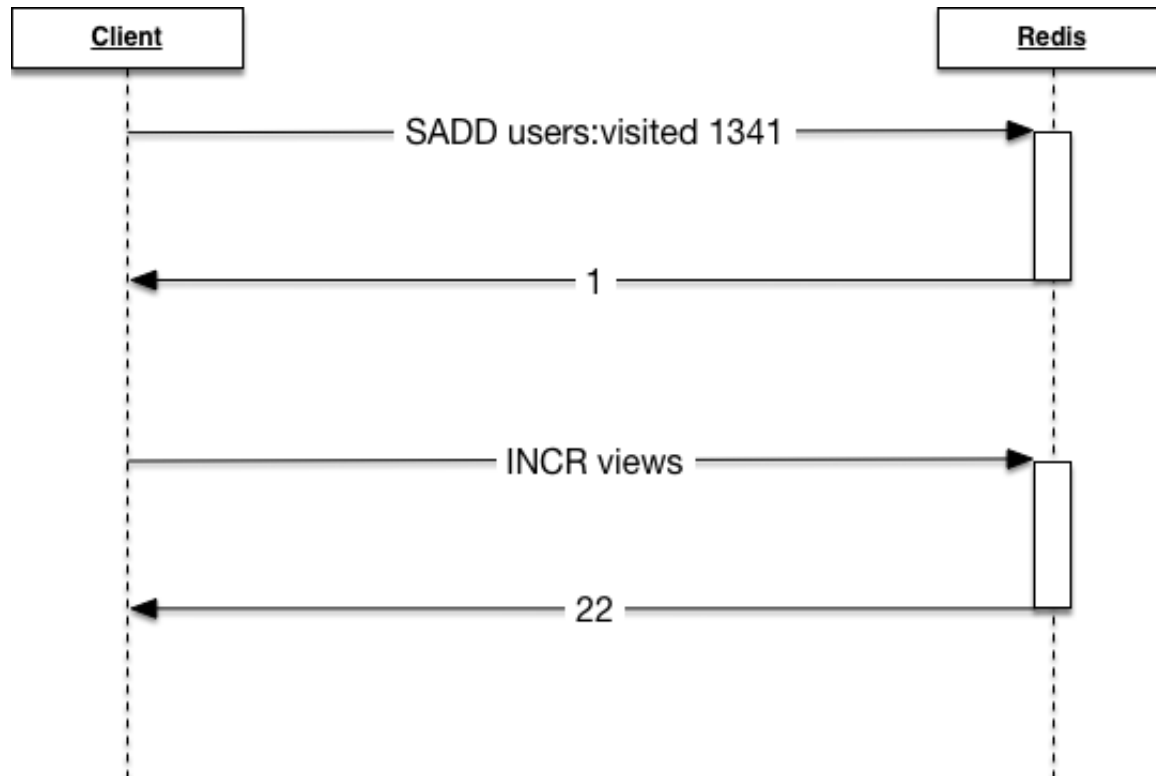
D – Durability

- Transaction changes are available event in the event of failure

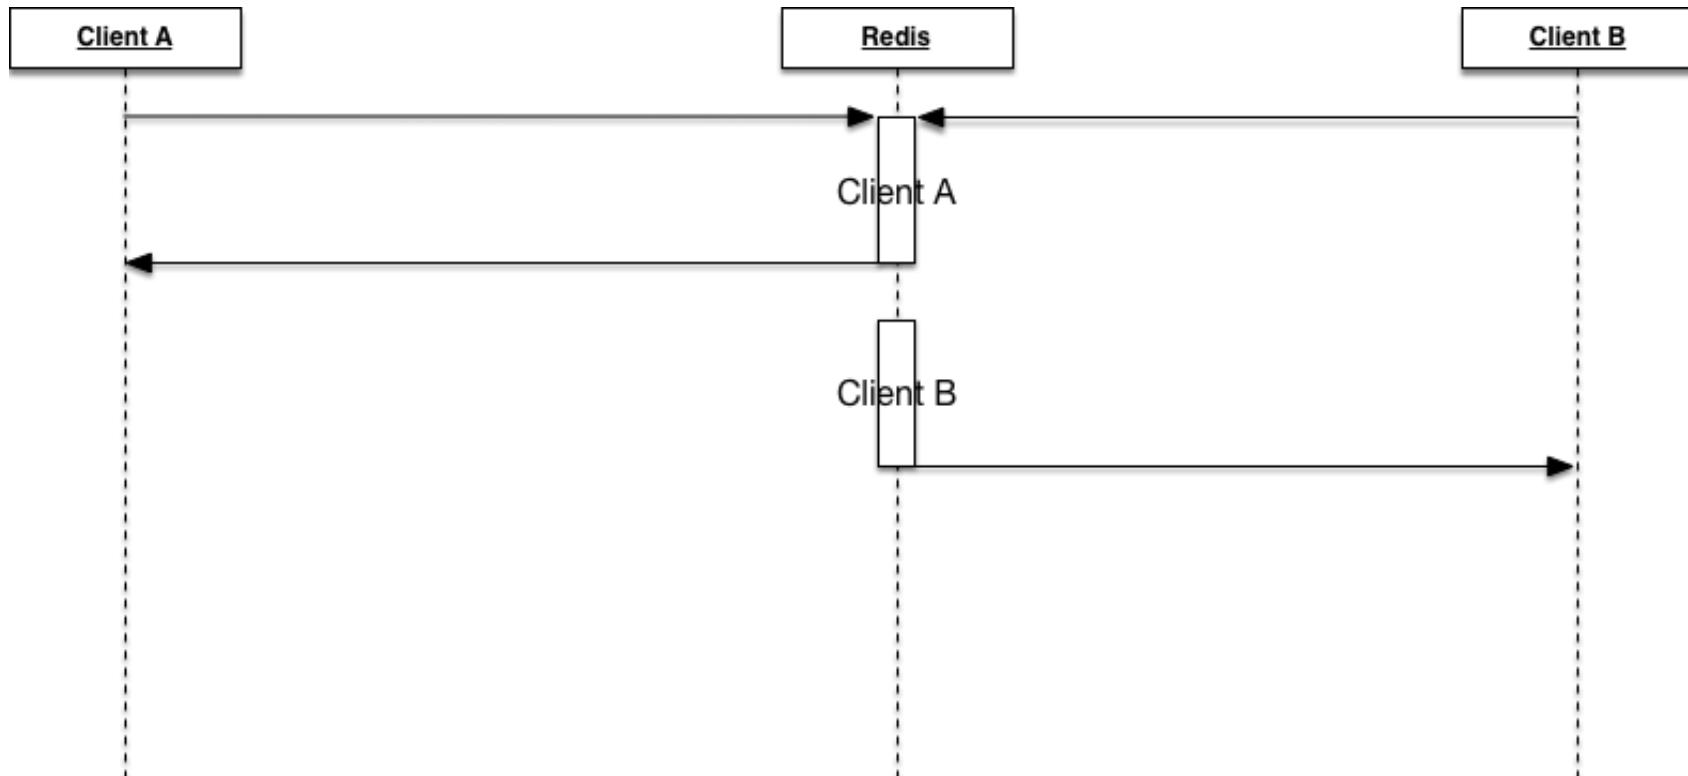
Redis Execution Model

redislabs

Single Client – Execution Flow



Two Client – Execution Flow



Blocking and Non-Blocking Commands

- Most Redis commands are synchronous
- Non-Blocking Commands
 - BGSAVE, BGREWRITEAOF
 - UNLINK (v4)
- Client Blocking Commands
 - SUBSCRIBE
 - BLPOP, BRPOP, BRPOPLPUSH
 - MONITOR
 - WAIT

Variadic (Dynamic Arity) Commands

- Variable number of arguments
- Examples (Non-exclusive)
 - MSET
 - MGET
 - HGET (v4)
 - HSET (v4)
- Executed as a single command

Multiple Command Transactions

redislabs

Multiple Command Transactions

- MULTI to start transaction block
- EXEC to close transaction block
- DISCARD to abort transaction block

- Commands are queued until exec
- All commands or no commands are applied
- Transactions can have errors

MULTI Example

```
127.0.0.1:6379> MULTI
OK
127.0.0.1:6379> sadd site:visitors 124
QUEUED
127.0.0.1:6379> incr site:raw-count
QUEUED
127.0.0.1:6379> hset sessions:124 userid tague ip 127.0.0.1
QUEUED
127.0.0.1:6379> EXEC
1) (integer) 1
2) (integer) 1
3) (integer) 2
```

DISCARD Example

```
127.0.0.1:6379> sadd site:visitors 124
QUEUED
127.0.0.1:6379> incr site:raw-count
QUEUED
127.0.0.1:6379> DISCARD
OK
```

Transactions with Errors – Syntactic Error

```
127.0.0.1:6379> MULTI
OK
127.0.0.1:6379> set site:visitors 10
QUEUED
127.0.0.1:6379> ste site:raw-count 20
(error) ERR unknown command 'ste'
127.0.0.1:6379> EXEC
(error) EXECABORT Transaction discarded because of previous errors.
```

Transactions with Errors – Semantic Error

```
127.0.0.1:6379> MULTI
OK
127.0.0.1:6379> set messages:hello "Hello World!"
QUEUED
127.0.0.1:6379> incr messages:hello
QUEUED
127.0.0.1:6379> EXEC
1) OK
2) (error) ERR value is not an integer or out of range
```

Conditional Execution/Optimistic Concurrency Control

- WATCH to conditionally execute transaction if key unchanged
 - UNWATCH clear c
 - DISCARD to aboard transaction block (CLI)
-
- All commands or no commands are applied
 - Transactions **can** have errors

Dependent Modifications– Incorrect Version

```
def incorrectCheckBalanceAndTransferAmount(debit, credit, amount):
    amount = float(amount)
    debitkey = debit
    creditkey = 'account:{}'.format(credit)
    fname = 'balance'
    balance = r.hget(debitkey, fname)

    # Potential race condition - start
    if balance >= amount:
        tx = r.pipeline()
        tx.hincrbyfloat(debitkey, fname, -amount)
        tx.hincrbyfloat(creditkey, fname, amount)
        return tx.execute()
    # Potential race condition - end
    else:
        raise Exception('insufficient funds')
```

Dependent Modifications – Correct Example

```
def checkBalanceAndTransferAmount(debit, credit, amount):
    amount = float(amount)
    debitkey = 'account:{}'.format(debit)
    creditkey = 'account:{}'.format(credit)
    fname = 'balance'
    while True:
        try:
            tx = r.pipeline()
            tx.watch(debitkey)
            balance = float(tx.hget(debitkey, fname))
            tx.multi()
            if balance >= amount:
                tx.hincrbyfloat(debitkey, fname, -amount)
                tx.hincrbyfloat(creditkey, fname, amount)
                return tx.execute()
            else:
                raise Exception('insufficient funds - time to get a job')
        except WatchError:
            # Reaching here means that the watched 'balance' value had changed,
            # so we can just retry or use any other backoff logic
            continue
```

Durability

redislabs

Disk Based Persistence

- Redis continues to serve commands from main memory
- Multiple Persistence modes
 - Snapshot Based (RDB)
 - Changelog based (AOF)
- Provides durability of data across power loss
 - Look into replication to prevent data loss in case of node loss

RDB Persistence

- Persistence
 - Fork Redis process
 - Child process writes new RDB file
 - Atomic replace old RDB file with new
- Configuration
 - SAVE directive (Redis.conf): SAVE <seconds> <min-changes>
 - Runtime : CONFIG SET SAVE "60 1000 120 100 180 1"
- Trigger manually
 - SAVEcommand (synch)
 - BGSAVE (backgroud)
- All commands or no commands are applied
- Transactions **can** have errors

AOF Persistence

- Configuration
 - APPENDONLY directive (Redis.conf): APPENDONLY YES
 - Runtime : CONFIG SET APPENDONLY YES
- AOF File fsynch options
 - Trade off speed for data security
 - Options: None, everysecond, always
- BGREWRITEAOF
 - aof file grows indefinitely
 - BGREWRITEAOF – trigger compaction of AOF file

Transaction Review

redislabs

Transactions

- *Mostly* ACID Transactions
 - Atomic – through MULTI/EXEC
 - Isolation, Consistency – single threaded nature
 - Durability – persistence modes
- No Rollback – transaction commands are queued then sent to server
- Single threaded event-loop for serving commands
- WATCH for optimistic concurrency control
- Pipelines for asynchronous commands

Review Questions

- What transaction guarantees does Redis provide?
- Give an example of how a Redis transaction can fail?
- What is the watch command used for?
- How to Redis transactions differ from SQL transactions?